USCC
Programming Assignment 3 – LLVM IR
By: Sanjay Madhav
University of Southern California

Useful Links

For this assignment, you will need to consult the LLVM documentation a great deal. There are three documents that will be especially helpful:

- [LLVM Language Reference Manual](#) – The "Instruction Reference" in particular will be useful to determine which instructions you will need to use for the variety of usage cases.
- [LLVM Programmers Manual](#) – There's a lot of information in this document, but make sure you read the class hierarchy section which describes the relationship between `User`, `Value`, and `Instruction`. The "Helpful Hints for Common Operations" section also has several examples of operations you will need to perform in PA3, PA4, and PA5.
- [IRBuilder Class Reference](#) – Throughout this assignment, you will use the `IRBuilder` to create instructions. So it will be helpful to keep this list of member functions handy.

Introduction

The endpoint of PA2 is a front-end for USCC that can identify syntactically and semantically valid USC programs. This generates an AST which can be printed to the console, but it cannot actually be executed in any meaningful way. In this assignment, you will *lower* the AST to LLVM IR. This lowering process will lose a great deal of information about the higher-level source program, but the benefit is that the LLVM IR is more easily optimized and can ultimately be converted to machine executable code.

By default, the USCC driver compiles into a `.bc` file which contains the LLVM IR (which is also referred to as LLVM *bitcode*). You can also instruct the driver to pretty-print the IR to stdout using the `-p` command line flag.

For instance, once you complete this assignment, you could use the following to compile and pretty-print the IR for the `ssa01.usc` program in the `tests` directory:

```
$ ../uscc/bin -p ssa01.usc
```

The output for the main function will look something along the lines of:

```
define i32 @main() {
entry:
  %y = alloca i32
  %x = alloca i32
  store i32 5, i32* %x
  store i32 6, i32* %y
  store i32 7, i32* %x
  store i32 10, i32* %x
  %x1 = load i32* %x
  %tobool = icmp ne i32 %x1, 0
  br i1 %tobool, label %if.then, label %if.else
```

```
if.then:                                              ; preds = %entry
  store i32 20, i32* %y
  br label %if.end
if.else:                                              ; preds = %entry
  store i32 15, i32* %y
  br label %if.end

if.end:                                               ; preds = %if.else, %if.then
  %y2 = load i32* %y
  %0 = call i32 (i8*, ...)* @printf(...)
  ret i32 0
}
```

The above example also illustrates the main properties of LLVM IR. Each function is represented with a *control flow graph* (CFG). A *basic block* in a CFG represents a sequential chain of execution, and it always ends with a *terminator instruction*, which represents a jump in control flow such as a branch or a return. Thus, the terminator instructions represent the edges in the CFG.

You'll also notice that LLVM IR looks very similar to assembly code, with the notable exception that there is an infinite number of registers called *virtual registers*. One key property of these registers is they can only be assigned to once – this is called *SSA form* which will become more relevant in PA4. For now, we will bypass SSA form by using the stack to store local variables.

The test suite associated with this programming assignment is testEmit.py. This test suite first compiles USC code into LLVM bitcode, and then interprets this bitcode via lli. The output from executing each program is then compared against the expected output. This ensures that the bitcode is both well-formed and functionally correct.

To run this test suite, execute the following command from the tests directory:

```
$ python testEmit.py
```

When you initially run this test suite, all 21 tests will fail. As you start filling in code in parse/ASTEmit.cpp, more and more tests will pass.

It is also possible to compile LLVM bitcode into platform-specific assembly via the llc tool. Once you have platform-specific assembly, it can be assembled into a native executable.

In this assignment, the vast majority of the code you will write will be in parse/ASTEmit.cpp, which contains the implementations of the emitIR function for each type of ASTNode. The emitIR function takes in a CodeContext by reference and returns a pointer to an llvm::Value. The CodeContext struct is defined in parse/Emitter.h. Of particular note is that the CodeContext contains a pointer to the current basic block called mBlock. This allows each emitIR call to know where the instructions should be added. This also means nodes that create basic blocks must update mBlock. The other member variables are used in some specific cases – for example mGlobal can be used to access the global LLVM context.

Generally, the recommended way to add instructions to a basic block is by using the IRBuilder class. You can declare an instance of IRBuilder as follows:

```
IRBuilder<> build(block);
```

Where `block` is the basic block you want to add instructions to. Once you create the `IRBuilder` instance, there are many functions prefixed with `Create` that can be used to append instructions to the basic block. Conveniently, these functions return a pointer to the instruction (which for most instructions corresponds to a `Value` pointer).

Implementation

It is recommended you implement the `emitIR` functions in the order outlined below. If you do not follow this order, it will be more difficult to pinpoint the nodes that are not implemented properly.

ASTExprStmt

`ASTExprStmt` is very straightforward -- it simply needs to call `emitIR` on the `mExpr` member variable. In this case, since an expression statement simply evaluates and does not return any value, the `emitIR` function should just return `nullptr`.

ASTCompoundStmt

A compound statement contains a list of declarations, followed by a list of statements. So this node just needs to emit each of the declarations, and then each of the statements. As with expression statements, we do not expect the node to return any values.

ASTConstantExpr

A constant expression is an expression representing a constant number. To create a `ConstantInt`, use the `ConstantInt::get` static method. This requires a type, which you can get using `Type::getInt32Ty` or `getInt8Ty`, depending on whether the constant is an integer or a character. Since this `ConstantInt` will be used, you must return the pointer to it.

ASTReturnStmt

This is the first node that will actually generate instructions via an `IRBuilder`. You will want to either `CreateRet` or `CreateRetVoid`, depending on the type of the statement. If the return is not void, you must first `emitIR` for the expression. This instruction does not need to be returned by `emitIR`.

At this point, three tests should pass: `emit02`, `emit05`, and `emit10`.

ASTBinaryMathOp

This node should generate the lhs and rhs expressions, and then create and return appropriate instruction based on the op: `add` (+), `sub` (-), `mul` (*), `sdiv` (/), and `srem` (%). Note that for instructions that return values, you can provide a name hint that corresponds to a suggested name for the virtual register that'll store the result. It's recommended you use names that make sense like "add", "sub", and so on. If the name is already in use, LLVM will automatically append a number to it.

ASTBinaryCmpOp

This is very similar to binary math ops, except you are generating `ICmp` instructions. One other difference is you need to zero-extend (`Zext`) the result of the `ICmp` math op into a `Int32Ty`. The reason for this is to allow for the C-style conversion of comparisons into integers.

At this point, the `emit03` test should also pass.

Creating Stack Space for Local Variables

For now, all local variables will exist on the stack. This is to work around the fact that LLVM uses an *SSA form*, meaning each virtual register can only be assigned once. In order to support this, you need to allocate stack space for these variables in `parse/Symbols.cpp` – specifically in `ScopeTable::emitIR`. Each function has a `ScopeTable` that encapsulates all of the variables declared within the function. Then in the entry block, all of these variables have space allocated.

The `emitIR` function is already provided for arrays, because arrays are somewhat complex. But for regular variables, you just need to create an `alloca` instruction. When you `alloca`, you will need to also supply the correct LLVM type, which you can get with the `llvmType` member function of identifier. The identifier's address should then be set to the result of the `alloca`.

There is one special case to handle, which is the case where the address already exists. This means that the local variable in question is a function parameter. In this case, the "address" actually contains the value of the parameter. So you still want to create an `alloca` instruction, but you also need to store the existing value into the newly allocated address. This creates a copy of the parameter (eg. pass by value).

Reading/Writing to Local Variables

Reading and writing of local variables is encapsulated by `Identifier::readFrom` and `Identifier::writeTo`. As with the previous part, the code for arrays is provided for you. You just need to create the `load` and `store` instructions for normal variables in the else cases. Remember that the address to load and store from is saved in the identifier.

At this `point`, `test015`, `emit04`, and `emit09` should also pass.

ASTIncExpr and ASTDecExpr

For these two nodes, you'll want to read the identifier using `readFrom`, add or subtract one to the value, and then write it back into the identifier with `writeTo`. Once these two nodes are implemented, `emit06` should also pass.

ASTNotExpr

For this node, you can't simply generate the sub expression and then `CreateNot`. This is because `CreateNot` just performs a one's complement. Instead, you want to `ICmpEQ` against zero and then `ZExt` to a 32-bit integer. This is because the not of anything other than zero should return zero.

ASTWhileStmt

This node will require creating basic blocks via `BasicBlock::Create`. The control flow of a while loop should be as follows:

This means the predecessor block has an unconditional branch to the `while.cond` block, which will either branch to `while.body` or `while.end`, depending on the result of the condition expression. The `while.body` block has an unconditional branch to the `while.cond` block. Note that when you create a basic block, you can pass in a suggested name for the block via the second parameter. It's recommend you follow the same naming convention as in the diagram.

The other aspect you need to make sure of is that you update the `mBlock` in the `CodeContext` as appropriate. So for example, when you are emitting the condition, you need to make sure `mBlock` is set to the `while.cond`.

Once while loops are implemented, you should now also pass the `emit07`, `emit08`, `emit09`, `opt05`, and `opt06` test cases.

ASTAssignStmt

This node simply needs to emit the expression and write the value of the expression to the identifier.

ASTIfStmt

If statements have two possibilities – either there is an if statement by itself, or an if statement with an associated else statement. The control flow of an if statement should be as follows:
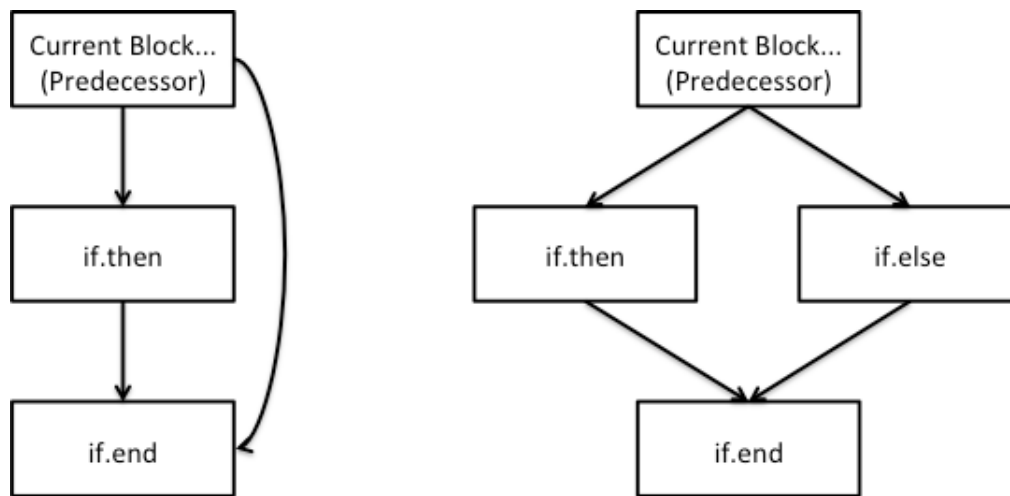
As with the while loop, you will need to create the appropriate basic blocks and generate the code in these basic blocks.

Once you implement if statements, your code should now pass all of the remaining test cases.

Conclusion

You now have a functional compiler in that all USC source programs will compile into LLVM bitcode. If you wanted to, you could also emit actual machine code for your machine via `llc`. The next two labs will be focused on improvements to the LLVM IR generation. In PA4, you will greatly reduce the number of stack memory instructions by implementing generation of SSA form for local variables. In PA5, you will implement optimization passes that will further improve execution time of the generated bitcode.